

225229140 pml lab 6

step 1

In [1]:

```
import pandas as pd
from sklearn.linear_model import LogisticRegression
```

In [2]:

```
data=pd.read_csv('diabetes.csv')
data
```

Out[2]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

In [3]:

```
data.head
```

Out[3]:

<bound method NDFrame.head of									Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6									
1	1	85	66	29	0	26.6									
2	8	183	64	0	0	23.3									
3	1	89	66	23	94	28.1									
4	0	137	40	35	168	43.1									
..									
763	10	101	76	48	180	32.9									
764	2	122	70	27	0	36.8									
765	5	121	72	23	112	26.2									
766	1	126	60	0	0	30.1									
767	1	93	70	31	0	30.4									

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]>

In [4]:

```
data.shape
```

Out[4]:

(768, 9)

In [5]:

```
data.columns
```

Out[5]:

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

In [6]:

```
data.dtypes
```

Out[6]:

```
Pregnancies      int64
Glucose           int64
BloodPressure     int64
SkinThickness     int64
Insulin           int64
BMI               float64
DiabetesPedigreeFunction float64
Age               int64
Outcome           int64
dtype: object
```

In [7]:

```
data.info
```

Out[7]:

```
<bound method DataFrame.info of      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0              6    148             72         35         0  33.6
1              1     85             66         29         0  26.6
2              8    183             64          0         0  23.3
3              1     89             66         23        94  28.1
4              0    137             40         35       168  43.1
..          ...    ...             ...         ...         ...   ...
763           10    101             76         48       180  32.9
764              2    122             70         27         0  36.8
765              5    121             72         23       112  26.2
766              1    126             60          0         0  30.1
767              1     93             70         31         0  30.4

      DiabetesPedigreeFunction  Age  Outcome
0              0.627         50         1
1              0.351         31         0
2              0.672         32         1
3              0.167         21         0
4              2.288         33         1
..          ...    ...             ...
763           0.171         63         0
764           0.340         27         0
765           0.245         30         0
766           0.349         47         1
767           0.315         23         0

[768 rows x 9 columns]>
```

In [8]:


```
data.Pregnancies.value_counts()
```

Out[8]:

```
1    135
0    111
2    103
3     75
4     68
5     57
6     50
7     45
8     38
9     28
10    24
11    11
13    10
12     9
14     2
15     1
17     1
Name: Pregnancies, dtype: int64
```

step 2

```
In [9]:
import seaborn as sns
sns.heatmap(data.head(10), annot=True)
```



step 3

```
In [10]:
X = data[['Age']]
y = data[['Outcome']]
```

```
In [11]:
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=42)
```

```
In [12]:
X_train
```



```
In [13]:
from sklearn import linear_model
logr=linear_model.LogisticRegression()
logr.fit(X,y)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
return f(*args, **kwargs)

Out[13]:

LogisticRegression()

In [14]:

```
df=logr.predict(X_test)
df
```

Out[14]:

```
array([0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0,
       0, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
       0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0], dtype=int64)
```

In [15]:

```
print("coef_ : ",logr.coef_)
print("intercept_ : ",logr.intercept_)
```

```
coef_ : [[0.04202466]]
intercept_ : [-2.04744865]
```

In [16]:

```
logr.predict([[60]])
```

Out[16]:

```
array([1], dtype=int64)
```

In [17]:

```
lrf=logr.coef_ * 60 + logr.intercept_
from scipy.special import expit
d=expit(lrf)
```

In [18]:

```
if d > 0.5:
    print('Yes, he will become diabetic ')
else:
    print('No, he will not be diabetic')
```

Yes, he will become diabetic

step 4

In [19]:

```
X1=data[['Glucose', 'BMI', 'Age']]
```

In [20]:

```
from sklearn import linear_model
```

```
X1_train,X1_test,y1_train,y1_test = train_test_split(X1,y,random_state=42,test_size=0.24)
logr1 =linear_model.LogisticRegression()
logr1.fit(X1_train,y1_train)
logr1.predict(X1_test)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

Out[20]:

```
array([0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0,
       0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0,
       0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1,
       0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0,
       0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0,
       0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1,
       0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
       0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0,
       0, 0, 0, 0, 1, 1, 0, 1, 1], dtype=int64)
```

In [21]:

```
print("coef_ : ",logr1.coef_)
print("intercept_ : ",logr1.intercept_)
```

```
coef_ : [[0.03292234 0.09635698 0.04398021]]
intercept_ : [-9.39683405]
```

In [22]:

```
lrf1=logr1.coef_ * 150 * 30 * 40+ logr1.intercept_
from scipy.special import expit
expit(lrf1)
```

Out[22]:

```
array([[1., 1., 1.]])
```

In [23]:

```
logr1.predict([[150,30,40]])
```

Out[23]:

```
array([1], dtype=int64)
```

In [24]:

```
logr1.predict_proba([[150,30,40]])
```

Out[24]:

```
array([[0.45228691, 0.54771309]])
```

step 5

In [25]:

```
X2=data.drop(['Outcome'],axis=1)
X2_train,X2_test,y2_train,y2_test = train_test_split(X2,y,test_size=.25,random_state=42)
from sklearn import linear_model
logr2=LogisticRegression()
logr2.fit(X2_train,y2_train)
df1=logr2.predict(X2_test)
df1
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Out[25]:

```
array([0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0,
       1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0,
       0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1,
       0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0,
       0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1,
       0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1,
       0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
       0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0,
       0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0])
```

In [26]:

```
from sklearn.metrics import roc_auc_score
lor_auc=roc_auc_score(y2_test,df1)
print("Auc:",lor_auc)
```

```
Auc: 0.7122658183103571
```

step 6

In [27]:

```
def get_auc(var,tar,df):
    fx = df[var]
    fy = df[tar]
    logr3=LogisticRegression()
    logr3.fit(fx,fy)
    pred=logr3.predict_proba(fx)[:,-1]
    auc_val = roc_auc_score(y,pred)
    return auc_val
get_auc(['Glucose','BMI'],['Outcome'],data)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

Out[27]:

0.8109328358208956

In [28]:

```
get_auc(['Pregnancies','BloodPressure','SkinThickness'],['Outcome'],data)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

Out[28]:

0.6444962686567164

In [29]:

```
def best_next(current,cand,tar,data):
    best_auc=-1
    best_var=None
    for i in cand:
        auc_v = get_auc(current+[i],tar,data)
        if auc_v>=best_auc:
            best_auc=auc_v
            best_var=i
    return best_var
```

In [30]:

```
current=['Insulin','BMI','DiabetesPedigreeFunction','Age']
cand=['Pregnancies','Glucose','BloodPressure','SkinThickness']
tar=['Outcome']
next_var = best_next(current,cand,tar,data)
next_var
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

Out[30]:

'Pregnancies'

In [31]:

```

tar = ['Outcome']
current=[]
cand=['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']
max_num = 7
num_it = min(max_num, len(cand))
for i in range(0, num_it):
    next_var = best_next(current, cand, tar, data)
    current += [next_var]
    cand.remove(next_var)
    print("variable added in step "+str(i+1)+" is "+ next_var + " .")

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

```
variable added in step 1 is Pregnancies .
```

```
variable added in step 2 is Glucose .
```

```
variable added in step 3 is BloodPressure .
```

```
variable added in step 4 is SkinThickness .
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
variable added in step 5 is Insulin .
```

```
variable added in step 6 is BMI .
```

```
variable added in step 7 is DiabetesPedigreeFunction .
```

In [32]:

```
print(current)
```

```
['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction']
```

step 7

In [33]:

```
X2_train, X2_test, y2_train, y2_test = train_test_split(X2, y, stratify=y, test_size=.5, random_state=42)
```

In [34]:

```
prediction = logit2.predict_proba(X2_test)
```

In [35]:

```

train = pd.concat([X2_train,y2_train],axis =1)
test = pd.concat([X2_test,y2_test],axis =1)
def auc_train_test (variables,target, train, test):
    X_train = train[variables]
    X_test = test[variables]
    Y_train =train[target]
    Y_test = test[target]
    Lor=LogisticRegression()
    Lor.fit(X_train,Y_train)
    prediction_train = Lor.predict_proba(X_train)[:,1]
    prediction_test = Lor.predict_proba(X_test)[:,1]
    auc_train = roc_auc_score(Y_train, prediction_train)
    auc_test = roc_auc_score(Y_train,prediction_test)
    return (auc_train, auc_test)
auc_values_train=[]
auc_values_test=[]
variable_evaluate=[]
for v in X2.columns:
    variable_evaluate.append(v)
    auc_train, auc_test = auc_train_test(variable_evaluate, ['Outcome'], train, test)
    auc_values_train.append(auc_train)
    auc_values_test.append(auc_test)

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
return f(*args, **kwargs)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:763: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)


```
n_iter_i = _check_optimize_result(
```

In []:

In []:

In [37]:


```
import matplotlib.pyplot as plt
import numpy as np
x = np.array(range(0, len(auc_values_train)))
my_train = np.array(auc_values_train)
my_test = np.array(auc_values_test)
plt.xticks(x, X2.columns, rotation=90)
plt.plot(x, my_train)
plt.plot(x, my_test)
plt.ylim(0.5, 1)
plt.show()
```




step 8

In [38]:

```
import scikitplot as skplt
skplt.metrics.plot_cumulative_gain(y2_test, prediction)
plt.show()
plt.figure(figsize=(7,7))
skplt.metrics.plot_lift_curve(y2_test, prediction)
plt.show()
```



<Figure size 504x504 with 0 Axes>



In []:

