

225229140 PML LAB 9

```
In [69]: import pandas as pd
```

```
In [70]: df=pd.read_csv('Employee_hopping.csv')
df
```

Out[70]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Ed
0	41	Yes	Travel_Rarely	1102	Sales	1	2	
1	49	No	Travel_Frequently	279	Research & Development	8	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	
4	27	No	Travel_Rarely	591	Research & Development	2	1	
5	32	No	Travel_Frequently	1005	Research & Development	2	2	
6	59	No	Travel_Rarely	1324	Research & Development	3	3	

```
In [71]: df.head()
```

Out[71]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educati
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life S
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life S
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life S
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns

```
In [72]: df.head
```

```
Out[72]: <bound method NDFrame.head of
te      Age Attrition      BusinessTravel  DailyRa
0      41      Yes      Travel_Rarely      1102      Sales
1      49      No      Travel_Frequently      279      Research & Development
2      37      Yes      Travel_Rarely      1373      Research & Development
3      33      No      Travel_Frequently      1392      Research & Development
4      27      No      Travel_Rarely      591      Research & Development
5      32      No      Travel_Frequently      1005      Research & Development
6      59      No      Travel_Rarely      1324      Research & Development
7      30      No      Travel_Rarely      1358      Research & Development
8      38      No      Travel_Frequently      216      Research & Development
9      36      No      Travel_Rarely      1299      Research & Development
10     35      No      Travel_Rarely      809      Research & Development
11     29      No      Travel_Rarely      153      Research & Development
12     31      No      Travel_Rarely      670      Research & Development
13     34      No      Travel_Rarely      1346      Research & Development
14     28      Yes      Travel_Rarely      103      Research & Development
15     29      No      Travel_Rarely      1389      Research & Development
16     32      No      Travel_Rarely      334      Research & Development
17     33      No      Travel_Rarely      1102      Research & Development
```

```
In [73]: df.shape
```

```
Out[73]: (1470, 35)
```

```
In [74]: df.columns
```

```
Out[74]: Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department',
'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount',
'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate',
'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction',
'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked',
'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating',
'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel',
'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance',
'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion',
'YearsWithCurrManager'],
dtype='object')
```

```
In [75]: df.dtypes
```

```
Out[75]: Age                int64
Attrition                 object
BusinessTravel            object
DailyRate                int64
Department               object
DistanceFromHome          int64
Education                 int64
EducationField            object
EmployeeCount             int64
EmployeeNumber            int64
EnvironmentSatisfaction   int64
Gender                   object
HourlyRate                int64
JobInvolvement            int64
JobLevel                  int64
JobRole                   object
JobSatisfaction           int64
MaritalStatus             object
MonthlyIncome             int64
MonthlyRate               int64
NumCompaniesWorked        int64
Over18                    object
OverTime                  object
PercentSalaryHike         int64
PerformanceRating         int64
RelationshipSatisfaction   int64
StandardHours             int64
StockOptionLevel          int64
TotalWorkingYears         int64
TrainingTimesLastYear     int64
WorkLifeBalance           int64
YearsAtCompany            int64
YearsInCurrentRole        int64
YearsSinceLastPromotion   int64
YearsWithCurrManager       int64
dtype: object
```

```
In [76]: df['Department'].value_counts()
```

```
Out[76]: Research & Development    961
Sales                             446
Human Resources                     63
Name: Department, dtype: int64
```

```
In [77]: #step 2
```

```
In [78]: x=df.drop(['Attrition'],axis=1)
y=df.Attrition
```

In [79]: x

22	34	Travel_Rarely	419	Research & Development	7	4	Life Science
23	21	Travel_Rarely	391	Research & Development	15	2	Life Science
24	34	Travel_Rarely	699	Research & Development	6	1	Medic
25	53	Travel_Rarely	1282	Research & Development	5	3	Oth
26	32	Travel_Frequently	1125	Research & Development	16	1	Life Science
27	42	Travel_Rarely	691	Sales	8	4	Marketi
28	44	Travel_Rarely	477	Research & Development	7	4	Medic

In [80]: y.head()

```
Out[80]: 0    Yes
1     No
2     Yes
3     No
4     No
Name: Attrition, dtype: object
```

In [81]: y=y.apply(lambda x:1 if x=='Yes' else 0)
y.head()

```
Out[81]: 0     1
1     0
2     1
3     0
4     0
Name: Attrition, dtype: int64
```

In [82]: #step 3

```
In [83]: df=pd.get_dummies(df,columns=['BusinessTravel','Department','EducationField','Gender'])
df.head()
```

Out[83]:

	Age	Attrition	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	Enrollment
0	41	Yes	1102	1	2	1	1	1
1	49	No	279	8	1	1	2	2
2	37	Yes	1373	2	2	1	4	4
3	33	No	1392	3	4	1	5	5
4	27	No	591	2	1	1	7	7

5 rows × 56 columns

```
In [84]: #step4
```

```
In [85]: X=df.drop(['Attrition'],axis=1)
```

```
In [86]: x.shape
```

Out[86]: (1470, 34)

```
In [87]: y.shape
```

Out[87]: (1470,)

```
In [88]: #step 5
```

```
In [92]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X,y, test_size =0.2, random_s
```

```
In [93]: from sklearn.ensemble import RandomForestClassifier
RFC = RandomForestClassifier(n_estimators=100, max_features=0.3)
```

```
In [94]: RFC.fit(x_train,y_train)
```

Out[94]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features=0.3, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
oob_score=False, random_state=None, verbose=0,
warm_start=False)


```
In [106]: feature_name = pd.DataFrame(RFC.feature_importances_, index=x_train.columns, columns=feature_name)
```


EducationField_Medical	0.004580
EducationField_Other	0.002800
EducationField_Technical Degree	0.007086
Gender_Female	0.006063
Gender_Male	0.005492
JobRole_Healthcare Representative	0.001381
JobRole_Human Resources	0.003253
JobRole_Laboratory Technician	0.008137
JobRole_Manager	0.000843
JobRole_Manufacturing Director	0.002220
JobRole_Research Director	0.000664
JobRole_Research Scientist	0.005193
JobRole_Sales Executive	0.007457

```
In [ ]:
```

```
In [110]: import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [114]: pd.Series(RFC.feature_importances_, index=x_train.columns).sort_values(ascending=
```

Out[114]: <matplotlib.axes._subplots.AxesSubplot at 0x1f129362cf8>

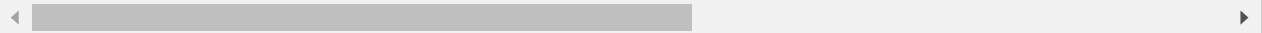


```
In [115]: #step 8
```



```
In [121]: estimator=RFC.estimators_[5]
```

```
In [125]: from sklearn import tree
from sklearn.tree import export_graphviz
with open("RFDT.dot", 'w') as f:
    f= tree.export_graphviz(estimator, out_file=f, max_depth=4, impurity=False, fe
```



```
In [126]: !dot-Tpng RFDT.dot -o RFDT.png
```

'dot-Tpng' is not recognized as an internal or external command,
operable program or batch file.

```
In [ ]:
```

```
In [130]: #step 9
```

```
In [138]: rf2 = RandomForestClassifier(oob_score=True, random_state=42, warm_start=True, n_
oob_list = list()
for n_trees in [15, 20, 30, 40, 50, 100, 150, 200, 300, 400]:
    rf2.set_params(n_estimators=n_trees)
    rf2.fit(x_train, y_train)
    oob_error = 1 - rf2.oob_score_
    oob_list.append(pd.Series({'n_trees': n_trees, 'oob': oob_error}))
rf_oob_df = pd.concat(oob_list, axis=1).T.set_index('n_trees')
rf_oob_df
```

C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\ensemble\forest.py:453: UserWarning: Some inputs do not have OOB scores. This probably means too few trees were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\ensemble\forest.py:458: RuntimeWarning: invalid value encountered in true_divide

predictions[k].sum(axis=1)[: , np.newaxis])

C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\ensemble\forest.py:453: UserWarning: Some inputs do not have OOB scores. This probably means too few trees were used to compute any reliable oob estimates.

warn("Some inputs do not have OOB scores. "

C:\Program Files (x86)\Microsoft Visual Studio\Shared\Anaconda3_64\lib\site-packages\sklearn\ensemble\forest.py:458: RuntimeWarning: invalid value encountered in true_divide

predictions[k].sum(axis=1)[: , np.newaxis])


Out[138]:

	oob
n_trees	
15.0	0.163265
20.0	0.159014
30.0	0.148810
40.0	0.144558
50.0	0.139456
100.0	0.140306
150.0	0.138605
200.0	0.140306
300.0	0.137755
400.0	0.140306

In [139]: *#step 10*

```
In [140]: ax = rf_oob_df.plot(legend=False, marker='o', figsize=(10,5))
ax.set(ylabel='out-of-bag error')
```

```
Out[140]: [Text(0,0.5,'out-of-bag error')]
```



```
In [141]: #step 11
```

```
In [143]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
clf = DecisionTreeClassifier(max_depth=4, random_state=42)
clf.fit(x_test, y_test)
```

```
Out[143]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=4,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=42,
splitter='best')
```



```
In [155]: print("RF model :",accuracy_score(y_test,RFC_y_pred))
print("RF Precision:",precision_score(y_test,RFC_y_pred))
print("RF Recall :",recall_score(y_test,RFC_y_pred))
print("RF F1 score :",f1_score(y_test,RFC_y_pred))
print("\n")
print("DT model :",accuracy_score(y_test,y_pred1))
print("DT Precision:",precision_score(y_test,y_pred1))
print("DT Recall :",recall_score(y_test,y_pred1))
print("DT F1 score :",f1_score(y_test,y_pred1))
```

```
RF model : 0.8741496598639455
RF Precision: 0.625
RF Recall : 0.1282051282051282
RF F1 score : 0.21276595744680848
```

```
DT model : 0.9183673469387755
DT Precision: 1.0
DT Recall : 0.38461538461538464
DT F1 score : 0.5555555555555556
```

```
In [ ]:
```