

SURIYA\_225229140

## Lab 10: Patients Physical Activities prediction using Boosting

### Step 1: Understand Data

```
In [6]: import pandas as pd
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
from sklearn.metrics import precision_score, recall_score, accuracy_score, roc_auc_score
from sklearn.ensemble import GradientBoostingClassifier, AdaBoostClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegressionCV
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score
```


```
In [7]: df = pd.read_csv("Human_Activity_Data.csv")
```

```
In [8]: df.head()
```

```
Out[8]:
```

	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAcc-mad()-X	tBodyAcc-mad()-Y
0	0.288585	-0.020294	-0.132905	-0.995279	-0.983111	-0.913526	-0.995112	-0.983185
1	0.278419	-0.016411	-0.123520	-0.998245	-0.975300	-0.960322	-0.998807	-0.974914
2	0.279653	-0.019467	-0.113462	-0.995380	-0.967187	-0.978944	-0.996520	-0.963668
3	0.279174	-0.026201	-0.123283	-0.996091	-0.983403	-0.990675	-0.997099	-0.982750
4	0.276629	-0.016570	-0.115362	-0.998139	-0.980817	-0.990482	-0.998321	-0.979672

5 rows × 562 columns



```
In [9]: df.shape
```

```
Out[9]: (151, 562)
```

In [10]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 151 entries, 0 to 150
Columns: 562 entries, tBodyAcc-mean()-X to Activity
dtypes: float64(561), object(1)
memory usage: 663.1+ KB
```

In [11]: `df.columns`

```
Out[11]: Index(['tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y', 'tBodyAcc-mean()-Z',
               'tBodyAcc-std()-X', 'tBodyAcc-std()-Y', 'tBodyAcc-std()-Z',
               'tBodyAcc-mad()-X', 'tBodyAcc-mad()-Y', 'tBodyAcc-mad()-Z',
               'tBodyAcc-max()-X',
               ...,
               'fBodyBodyGyroJerkMag-skewness()', 'fBodyBodyGyroJerkMag-kurtosis()',
               'angle(tBodyAccMean,gravity)', 'angle(tBodyAccJerkMean,gravityMea
n)',
               'angle(tBodyGyroMean,gravityMean)',
               'angle(tBodyGyroJerkMean,gravityMean)', 'angle(X,gravityMean)',
               'angle(Y,gravityMean)', 'angle(Z,gravityMean)', 'Activity'],
              dtype='object', length=562)
```

In [12]: `type(df)`

Out[12]: `pandas.core.frame.DataFrame`

In [13]: `df['Activity'].value_counts`

```
Out[13]: <bound method IndexOpsMixin.value_counts of 0          STANDING
1          STANDING
2          STANDING
3          STANDING
4          STANDING
...
146  WALKING_DOWNSTAIRS
147  WALKING_DOWNSTAIRS
148  WALKING_DOWNSTAIRS
149  WALKING_DOWNSTAIRS
150          NaN
Name: Activity, Length: 151, dtype: object>
```

## Step 2: Build a small dataset

```
In [14]: lay = df.loc[df['Activity'] == "LAYING"][:500]
sit = df.loc[df['Activity'] == "SITTING"][:500]
walk = df.loc[df['Activity'] == "WALKING"][:500]
frames = [lay, sit, walk]
df_new = pd.concat(frames)
```

```
In [15]: df_new.shape
```

```
Out[15]: (98, 562)
```

```
In [16]: df_new.to_csv("Human_Activity_sample.csv")
```


```
In [17]: df1=pd.read_csv('Human_Activity_sample.csv')
```

```
In [18]: df1.head()
```

```
Out[18]:
```

	Unnamed: 0	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-std()-X	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAcc-mad()-X
0	51	0.403474	-0.015074	-0.118167	-0.914811	-0.895231	-0.891748	-0.917696
1	52	0.278373	-0.020561	-0.096825	-0.984883	-0.991118	-0.982112	-0.987985
2	53	0.276555	-0.017869	-0.107621	-0.994195	-0.996372	-0.995615	-0.994901
3	54	0.279575	-0.017276	-0.109481	-0.996135	-0.995812	-0.998689	-0.996393
4	55	0.276527	-0.016819	-0.107983	-0.996775	-0.997256	-0.995422	-0.997167

5 rows × 563 columns



```
In [19]: df1.shape
```

```
Out[19]: (98, 563)
```

```
In [20]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 98 entries, 0 to 97
Columns: 563 entries, Unnamed: 0 to Activity
dtypes: float64(561), int64(1), object(1)
memory usage: 431.2+ KB
```

In [21]: `df1.columns`

Out[21]: Index(['Unnamed: 0', 'tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y', 'tBodyAcc-mean()-Z', 'tBodyAcc-std()-X', 'tBodyAcc-std()-Y', 'tBodyAcc-std()-Z', 'tBodyAcc-mad()-X', 'tBodyAcc-mad()-Y', 'tBodyAcc-mad()-Z', ..., 'fBodyBodyGyroJerkMag-skewness()', 'fBodyBodyGyroJerkMag-kurtosis()', 'angle(tBodyAccMean,gravity)', 'angle(tBodyAccJerkMean,gravityMean)', 'angle(tBodyGyroMean,gravityMean)', 'angle(tBodyGyroJerkMean,gravityMean)', 'angle(X,gravityMean)', 'angle(Y,gravityMean)', 'angle(Z,gravityMean)', 'Activity'], dtype='object', length=563)

In [22]: `type(df1)`

Out[22]: `pandas.core.frame.DataFrame`

In [23]: `df1["Activity"].value_counts()`

Out[23]: WALKING 47  
LAYING 27  
SITTING 24  
Name: Activity, dtype: int64

### Step 3: Build GradientBoostingClassifier

In [24]: `X=df1.drop('Activity',axis=1)`  
`y=df1.Activity`

In [25]: `X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=42)`

In [26]: `model = GradientBoostingClassifier(n_estimators=100,learning_rate=1.0,max_depth=1)`  
`model.fit(X_train,y_train)`

Out[26]: GradientBoostingClassifier(learning\_rate=1.0, max\_depth=1, random\_state=42)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [27]: `y_pred=model.predict(X_test)`

```
In [28]: accuracy_score(y_test,y_pred)
```

```
Out[28]: 0.9666666666666667
```

```
In [29]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
LAYING	1.00	0.90	0.95	10
SITTING	0.89	1.00	0.94	8
WALKING	1.00	1.00	1.00	12
accuracy			0.97	30
macro avg	0.96	0.97	0.96	30
weighted avg	0.97	0.97	0.97	30

#### Step4. [Find Best no. of trees and Best Learning Rate using Grid Search and Cross Validation]

```
In [30]: classifier = GradientBoostingClassifier()
```

```
In [31]: all_scores = cross_val_score(estimator=classifier, X=X_train, y=y_train, cv=5)
```

```
In [32]: all_scores
```

```
Out[32]: array([1., 1., 1., 1., 1.])
```

**To find the average of all the accuracies, simple use the mean() method**

```
In [33]: all_scores.mean()
```

```
Out[33]: 1.0
```

```
In [34]: parameter = {'n_estimators': [50, 100, 200, 400], 'learning_rate': [0.1, 0.01]}
```

```
In [35]: model1 = GridSearchCV(estimator=classifier, param_grid=parameter, cv=5, n_jobs=
```

```
In [36]: model1.fit(X_train,y_train)
```

```
Out[36]: GridSearchCV(cv=5, estimator=GradientBoostingClassifier(), n_jobs=-1,
                  param_grid={'learning_rate': [0.1, 0.01],
                              'n_estimators': [50, 100, 200, 400]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [41]: y_pred2=model1.predict(X_test)
```

```
In [42]: accuracy_score(y_test,y_pred2)
```

```
Out[42]: 1.0
```

```
In [43]: print(classification_report(y_test,y_pred2))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	10
SITTING	1.00	1.00	1.00	8
WALKING	1.00	1.00	1.00	12
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [44]: print(model1.best_estimator_)
```

```
GradientBoostingClassifier(n_estimators=50)
```

## Step5. [Build AdaBoostClassifier]

```
In [45]: base = DecisionTreeClassifier()
```

```
In [46]: model2 = AdaBoostClassifier(base_estimator=base,random_state=0)
```

```
In [47]: param_grid = {'n_estimators': [100, 150, 200], 'learning_rate': [0.01, 0.001]}
```

```
In [48]: model3 = GridSearchCV(model2,param_grid,cv=5,n_jobs=-1)
```

```
In [49]: model3.fit(X_train,y_train)
```

```
Out[49]: GridSearchCV(cv=5,
                      estimator=AdaBoostClassifier(base_estimator=DecisionTreeClassifier(),
                                                    random_state=0),
                      n_jobs=-1,
                      param_grid={'learning_rate': [0.01, 0.001],
                                   'n_estimators': [100, 150, 200]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [50]: y_pred3=model3.predict(X_test)
```

```
In [51]: accuracy_score(y_test,y_pred3)
```

```
Out[51]: 1.0
```

```
In [52]: print(classification_report(y_test,y_pred3))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	10
SITTING	1.00	1.00	1.00	8
WALKING	1.00	1.00	1.00	12
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [53]: print(model3.best_estimator_)
```

```
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), learning_rate=0.01,
                   n_estimators=100, random_state=0)
```

## Step6. [Build LogisticRegressionCV classifier]

```
In [54]: model4 = LogisticRegressionCV(cv=4,Cs=5,penalty='l2')
```

```
In [55]: model4.fit(X_train,y_train)
```

```
Out[55]: LogisticRegressionCV(Cs=5, cv=4)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [56]: y_pred4=model4.predict(X_test)
```

```
In [57]: accuracy_score(y_test,y_pred4)
```

```
Out[57]: 1.0
```

```
In [58]: print(classification_report(y_test,y_pred4))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	10
SITTING	1.00	1.00	1.00	8
WALKING	1.00	1.00	1.00	12
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

## Step 7 [ Build VotingClassifier]

```
In [59]: model5=VotingClassifier(estimators=[('lr',model4),('gbc',model1)], voting='ha
```



```
In [60]: model5.fit(X_train,y_train)
```

```
Out[60]: VotingClassifier(estimators=[('lr', LogisticRegressionCV(Cs=5, cv=4)),
                                      ('gbc',
                                       GridSearchCV(cv=5,
                                                    estimator=GradientBoostingClassif
ier(),
                                                    n_jobs=-1,
                                                    param_grid={'learning_rate': [0.
1,
                                                                0.0
1],
                                                                'n_estimators': [50,
100,
                                                                200,
                                                                40
0]}))])
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [63]: y_pred5=model5.predict(X_test)
```

```
In [64]: print(classification_report(y_test,y_pred5))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	10
SITTING	1.00	1.00	1.00	8
WALKING	1.00	1.00	1.00	12
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

## Step8. [ Interpret your results]

```
In [65]: print(model1.best_estimator_)
```

```
GradientBoostingClassifier(n_estimators=50)
```

```
In [66]: print(model3.best_estimator_)
```

```
AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), learning_rate=0.01,
                   n_estimators=100, random_state=0)
```

## GradientBoostingClassifier

### GradientBoostingClassifier(n\_estimators=50)

```
In [67]: classifierF = GradientBoostingClassifier(n_estimators=50)
all_scoresF = cross_val_score(estimator=classifier, X=X_train, y=y_train, cv=
parameter = {'n_estimators': [50, 100, 200, 400], 'learning_rate': [0.1, 0.01]
```

```
In [68]: modelGC = GridSearchCV(estimator=classifier, param_grid=parameter, cv=5, n_job
```

```
In [69]: modelGC.fit(X_train, y_train)
```

```
Out[69]: GridSearchCV(cv=5, estimator=GradientBoostingClassifier(), n_jobs=-1,
                    param_grid={'learning_rate': [0.1, 0.01],
                                'n_estimators': [50, 100, 200, 400]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [71]: y_predGC=model3.predict(X_test)
```

```
In [72]: accuracy_score(y_test, y_predGC)
```

```
Out[72]: 1.0
```

```
In [73]: print(classification_report(y_test, y_predGC))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	10
SITTING	1.00	1.00	1.00	8
WALKING	1.00	1.00	1.00	12
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
In [74]: ### AdaBoostClassifier
```

```
In [75]: #### AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), Learning_rate=0.01)
```

```
In [77]: modelABC = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), learning_rate=0.01)
```

```
In [78]: param_grid = {'n_estimators': [100, 150, 200], 'learning_rate': [0.01, 0.001]}
```

```
In [79]: modelGSCV = GridSearchCV(modelABC,param_grid,cv=5,n_jobs=-1)
modelGSCV.fit(X_train,y_train)
```

```
Out[79]: GridSearchCV(cv=5,
                    estimator=AdaBoostClassifier(base_estimator=DecisionTreeClassifier(),
                                                  learning_rate=0.01),
                    n_jobs=-1,
                    param_grid={'learning_rate': [0.01, 0.001],
                                'n_estimators': [100, 150, 200]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [80]: y_predGSCV=model3.predict(X_test)
```

```
In [81]: accuracy_score(y_test,y_predGSCV)
```

```
Out[81]: 1.0
```

```
In [82]: print(classification_report(y_test,y_predGSCV))
```

	precision	recall	f1-score	support
LAYING	1.00	1.00	1.00	10
SITTING	1.00	1.00	1.00	8
WALKING	1.00	1.00	1.00	12
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

